

AI.COREBLOCK zkTelemetry Proof Submissions via Kurier (zkVerify)

Author: Serkan Sahin (Development)

Project: AI.COREBLOCK - zkTelemetry → Kurier → zkVerify Mainnet

Document: Milestone 2 Technical Documentation

Date: 2026-02-08

1. Overview

Milestone 2 delivers a reliable, high-throughput proof submission pipeline from AI.COREBLOCK's zkTelemetry system to **zkVerify mainnet** via the **Kurier** submission layer. The milestone focuses on two core requirements:

1. **Durable job persistence:** Every proof submission is recorded and can be resumed or reconciled without losing correlation to Kurier jobs or on-chain transactions.
2. **Traceable on-chain verification:** For each submitted proof, we persist the zkVerify **extrinsic hash** and a direct **explorer URL**, enabling auditability and independent verification.

This milestone also documents the practical engineering process: the pipeline did not work perfectly from the start; we iterated through failures and edge cases and hardened the worker until it became stable at scale.

2. Architecture & Data Flow

2.1 Components

- **Supabase (PostgreSQL):** stores telemetry sessions and proof job state.
- **Proof Submission Worker (Node.js / TypeScript):** pulls queued proof jobs, submits them to Kurier, and persists results.
- **Kurier API:** job submission + job status tracking (finalization + tx hash).
- **zkVerify mainnet / Subscan explorer:** public on-chain record of proof submissions.

2.2 End-to-End Flow

1. Telemetry data is ingested from smartphones and stored in Supabase.
2. Proof jobs are created in Supabase with required proof inputs (e.g., commitment hash, proof data).
3. Worker pulls jobs in batches and submits them to Kurier.
4. Kurier returns a job identifier (jobId / requestId).
5. Once final, Kurier provides the zkVerify extrinsic hash and explorer link.
6. A sync script backfills tx hash / explorer URL for submitted jobs.

2.3 Smartphone Fleet in Vehicles

During Milestone 2 we operated a real-device pilot using **14 smartphones installed in vehicles** to collect telemetry under real driving conditions. This fleet was used to generate the input data that feeds the proof pipeline and to validate stable end-to-end execution at higher volume.

Milestone 2 Documentation

Final Version

Figure 1. Example view of the telemetry dataset captured from the smartphone fleet (Supabase table view).

This setup ensured that the system was tested against a broader range of real-world conditions (e.g., different routes and driving patterns) while keeping the proof submission flow fully traceable from Supabase job records to on-chain zkVerify extrinsics.


3. Supabase Schema & Job State Model

3.1 Core Tables (Milestone 2 scope)

- `smartphone_telemetry_sessions`
 - `telemetry_metrics`
 - `proof_jobs`
 - (optional views) `v_proof_jobs_ready`, `v_ai_coreblock_proofs`, etc.

Milestone 2 Documentation

Final Version




The screenshot shows the Supabase schema overview. On the left is a sidebar with icons for tables, views, and other database objects. The main area is a list of tables and views, each with a 'UNRESTRICTED' access level. The tables listed are: proof_jobs, scaled_smartphone_telemetry_sessions, smartphone_telemetry_sessions, telemetry_metrics, telemetry_proof_keys, v_ai_coreblock_proofs, v_aicb_proof_registry, and v_proof_jobs_ready.

Figure 2. Supabase schema overview showing telemetry and proof-related tables/views.

3.2 proof_jobs (Key Columns)

Each proof job is tracked with the following essential fields:

- status — lifecycle state (queued, submitting, submitted, verified, failed)
- attempts, last_error — resilience & debugging
- commitment_hash — unique proof reference (input-level identity)
- kurier_request_id — Kurier job identifier (critical for correlation)
- kurier_response — raw response payload snapshot for debugging/recovery
- zkverify_tx_hash — on-chain extrinsic hash
- zkverify_explorer_url — direct explorer link for verification
- created_at, updated_at, submitted_at, verified_at



The screenshot shows the Supabase Table Editor for the proof_jobs table. The table has columns: id (uuid), verified_at (timestamp), zkverify_tx_hash (text), zkverify_explorer_url (text), and created_at (timestamp). The table contains 29,424 records. The zkverify_tx_hash column shows various hex values, and the zkverify_explorer_url column shows URLs starting with <https://zkverify.subscan.io/extrinsic/>. The table is sorted by verified_at.

Figure 3. Supabase proof_jobs table showing persisted zkverify_tx_hash and zkverify_explorer_url.

3.3 Job Status Semantics

- **queued**: ready to be submitted
- **submitting**: reserved by worker, submission in progress
- **submitted**: Kurier accepted job; awaiting on-chain finalization or tx enrichment
- **verified**: job confirmed verified/finalized (if applicable in the flow)
- **failed**: permanent failure after retries or non-retryable error

4. Worker: Submission, Persistence & Reliability

4.1 Worker Responsibilities

The Proof Submission Worker is responsible for:

- fetching queued jobs in batches
- moving them to submitting
- submitting proofs to Kurier
- extracting and persisting:
 - kurier_request_id
 - raw response into kurier_response
 - (if available immediately) zkverify_tx_hash and zkverify_explorer_url
- updating final state (submitted / verified / failed) with retries and backoff logic




Figure 4. Project repository structure showing worker + Kurier sync scripts.

4.2 Running the Worker (High Throughput Mode)

To run the worker directly from TypeScript (avoiding ESM build resolution issues), we start it via tsx: BATCH_SIZE=100 CONCURRENCY=20 npm run dev

This starts the worker and continuously processes proof jobs from Supabase.

4.3 Configuration

The system is configured through environment variables (.env) including:

- SUPABASE_URL
- SUPABASE_SERVICE_ROLE_KEY
- KURIER_API_URL
- KURIER_API_KEY
- optional worker parameters:
 - BATCH_SIZE
 - CONCURRENCY
 - MAX_ATTEMPTS
 - POLL_INTERVAL_MS

```
⚙ .env
1 # Kurier API Configuration
2 KURIER_API_KEY=████████████████
3 KURIER_API_URL=https://api.kurier.xyz/api/
4
5 # Supabase Configuration
6 SUPABASE_URL=████████████████
7 SUPABASE_SERVICE_ROLE_KEY=ε████████████████
8
9 # Worker Configuration (optional - default
10 BATCH_SIZE=50
11 CONCURRENCY=2
12 MAX_ATTEMPTS=5
13 POLL_INTERVAL_MS=1500
```

Figure 5. Example .env configuration (API keys redacted).

4.4 Throughput & Load Testing

We intentionally tested the worker across a wide operational range—from near real-time confirmations (single-device submissions where proofs appear almost immediately after a smartphone session completes) up to **large batch submissions (e.g., ~1,000 proofs at once)** to stress-test end-to-end throughput.

This testing was essential: it helped us identify practical limits and bottlenecks across Kurier and zkVerify (rate limits, delayed transaction metadata availability, and optimal concurrency/batch sizing). Based on these findings, we calibrated the system to operate within a stable “sweet spot” that balances **speed, volume, and reliability**, while maintaining full traceability from Supabase job records to on-chain zkVerify extrinsics.

5. Kurier Job Tracking & tx Hash Synchronization

5.1 Why Sync Is Needed

Milestone 2 Documentation

Final Version

In high-volume operation, Kurier job submission can return quickly while the on-chain transaction hash becomes available slightly later. To ensure complete auditability, we synchronize and backfill missing on-chain fields for submitted jobs.

5.2 Sync Script

The milestone includes a sync script:

- zk/scripts/sync-kurier-status.mjs

It identifies submitted jobs where zkverify_tx_hash is missing but kurier_request_id exists, fetches the status from Kurier, and updates the database accordingly.

Because the script is typically run from a shell, .env needs to be loaded:

```
set -a; source .env; set +a; node zk/scripts/sync-kurier-status.mjs
```

6. On-chain Verification (zkVerify Explorer)

6.1 Explorer Evidence (Extrinsic Details)

Each submitted proof is verifiable on zkVerify mainnet via the extrinsic hash. The system stores both:

- zkverify_tx_hash
- zkverify_explorer_url

This allows direct linking from Supabase proof job to the public on-chain transaction.

Extrinsic #2088245-15	
Timestamp	2026-02-08 12:35:48 (UTC)
Block Time	34 secs ago
Block	2088245
Life Time	2088241 - 2088305
Extrinsic Hash	0xb7e9db3c4ade255a5a4d12d0b50ff8646540f467ded2a80380aa068b042d27
Action	SettlementGroth16pallet (Submit_proof)
Sender	ZKXwbGP2QyLv2F7JUrXKksKhTec39ytcGBNyTLfTZWucMmM8
Estimated Fee	0.017288156415 VFY
Used Fee	0.017288156415 VFY
Nonce	1347926
Result	Success

Figure 6. zkVerify Subscan extrinsic details showing Submit_proof with success status.

6.2 High-volume Evidence (Account Activity)

At scale, proof submissions appear as high-frequency extrinsics (e.g., submit_proof) on the explorer.

Milestone 2 Documentation

Final Version




Figure 7. Subscan account page showing high-volume extrinsics and proof submission activity.

7. Engineering Iteration & Hardening (What We Learned)

Engineering Note

Milestone 2 was achieved through iterative hardening. Early job runs revealed edge cases (e.g., missing correlation fields and delayed transaction hash availability at high throughput). These findings directly informed the final resilience measures (robust response parsing, strict persistence checks, and a reconciliation/sync path), resulting in stable high-volume operation.

We addressed these by improving:

- **robust response parsing** (extract job ID from multiple response shapes)
- **persistence guarantees** (ensuring Kurier job ID is stored)
- **sync & reconciliation paths** (ability to recover missing transaction metadata)
- **runbook clarity** (standardized worker and sync commands)

Milestone 2 Documentation

Final Version

Jobs

Refresh

↓	Updated At ↑	Job ID	Status	Chain	Transaction Hash	Block
7	2025-12-26 13:07:07	c7f4b346-e25b-11f0-99a3-e2579a7a7dd2	✖ Failed	-	-	-
	2025-12-26 13:01:11	f34752e4-e25a-11f0-99a3-e2579a7a7dd2	✖ Failed	-	-	-
3	2025-12-26 12:53:23	dc7fb1a1-e259-11f0-99a3-e2579a7a7dd2	✖ Failed	-	-	-
.9	2025-12-26 12:48:49	39636dc4-e259-11f0-99a3-e2579a7a7dd2	✖ Failed	-	-	-
2	2025-12-26 12:48:32	2f7113f9-e259-11f0-99a3-e2579a7a7dd2	✖ Failed	-	-	-
11	2025-12-26 12:48:01	1cf54ca0-e259-11f0-99a3-e2579a7a7dd2	✖ Failed	-	-	-
3	2025-12-26 12:15:39	9770ee59-e254-11f0-99a3-e2579a7a7dd2	✖ Failed	-	-	-
5	2025-12-26 12:10:45	e809568d-e253-11f0-99a3-e2579a7a7dd2	✖ Failed	-	-	-
5	2025-12-25 21:27:35	87706cc6-e1d8-11f0-99a3-e2579a7a7dd2	✖ Failed	-	-	-
7	2025-12-25 21:24:27	173faa6a-e1d8-11f0-99a3-e2579a7a7dd2	✖ Failed	-	-	-
2	2025-12-25 21:19:32	674ec29d-e1d7-11f0-99a3-e2579a7a7dd2	✖ Failed	-	-	-
7	2025-12-25 21:15:10	1f6f71a1-e1d6-11f0-99a3-e2579a7a7dd2	✖ Failed	-	-	-

Figure 8. Kurier “failed jobs” view from early iterations (before hardening).

Jobs

Refresh

Created At ↓	Updated At ↑	Job ID	Status	Chain	Transaction Hash
2026-02-08 00:21:32	2026-02-08 00:21:32	1ec08d6b-0484-11f1-99a3-e2579a7a7dd2	Finalized	-	0x20b7bb0d..12b3bfe3
2026-02-08 00:21:31	2026-02-08 00:21:31	1dbe2fec-0484-11f1-99a3-e2579a7a7dd2	Finalized	-	0xa4bc2787..cb9214252
2026-02-08 00:21:29	2026-02-08 00:21:29	1cba348b-0484-11f1-99a3-e2579a7a7dd2	Finalized	-	0x3287eb84..07bf84266
2026-02-08 00:21:27	2026-02-08 00:21:27	1ba802e9-0484-11f1-99a3-e2579a7a7dd2	Finalized	-	0x376fb78..5557a3391
2026-02-08 00:21:25	2026-02-08 00:21:25	1a865caf-0484-11f1-99a3-e2579a7a7dd2	Finalized	-	0xae45948f..57c3657e9
2026-02-08 00:21:24	2026-02-08 00:21:24	19800df5-0484-11f1-99a3-e2579a7a7dd2	Finalized	-	0x4536a183..fe94e35d7
2026-02-08 00:21:22	2026-02-08 00:21:22	1865d9d8-0484-11f1-99a3-e2579a7a7dd2	Finalized	-	0x93774e61..d28077a97
2026-02-08 00:21:20	2026-02-08 00:21:20	174c54ea-0484-11f1-99a3-e2579a7a7dd2	Finalized	-	0x549872c0..ab6820091
2026-02-08 00:21:18	2026-02-08 00:21:18	1646c447-0484-11f1-99a3-e2579a7a7dd2	Finalized	-	0xd425c785..c62faa85f
2026-02-08 00:21:16	2026-02-08 00:21:16	15498cb6-0484-11f1-99a3-e2579a7a7dd2	Finalized	-	0x14268144..c4cd440cf
2026-02-08 00:21:15	2026-02-08 00:21:15	143c3c43-0484-11f1-99a3-e2579a7a7dd2	Finalized	-	0x340660f2..5e2fb2a6c
2026-02-08 00:21:13	2026-02-08 00:21:13	132822c9-0484-11f1-99a3-e2579a7a7dd2	Finalized	-	0x9a3388bc..d94107dd8

Figure 9. Kurier “finalized jobs” view after stabilization, showing transaction hashes.

8. Note on Attribution: Kurier Sender Accounts

Kurier submits proofs using its own sender infrastructure. Therefore, it is not always reliable to filter on-chain submissions by sender address alone when multiple clients share the same sender accounts. **In practice, we attribute proofs using the per-job kurier_request_id and the corresponding zkverify_tx_hash persisted in Supabase.**

Milestone 2 Documentation

Final Version

Each AI.COREBLOCK proof is uniquely attributable via:

- kurier_request_id stored in Supabase for every proof job
- zkverify_tx_hash and zkverify_explorer_url backfilled via the sync script
- commitment_hash as the proof-level identifier in our system

Improvement Request

To make filtering easier at high volume, we propose attaching a project-specific identifier (e.g., verification key hash / vkHash, commitment_hash, or a clientReference tag) to each Kurier submission so Kurier-side filtering becomes trivial even when shared senders are used.

9. Conclusion

Milestone 2 successfully establishes a robust, high-throughput proof submission pipeline from AI.COREBLOCK's zkTelemetry system to zkVerify mainnet via Kurier. The solution provides end-to-end traceability by persisting correlation identifiers (kurier_request_id) and on-chain references (zkverify_tx_hash, zkverify_explorer_url) per proof job in Supabase, enabling independent verification through the zkVerify explorer.

In addition to achieving stable operation at higher volumes, the milestone resulted in a clear operational runbook (worker execution, monitoring queries, and transaction backfill via the sync script). The system is therefore ready to scale proof generation further.